

Amir Mirkazemi

Las Vegas, NV | amirkazemi@ucsd.edu | [linkedin.com/in/amir-mirkazemi](https://www.linkedin.com/in/amir-mirkazemi) | github.com/Amir-Mirkazemi | amirdata.com

EXPERIENCE

Data Engineer — Credit One Bank

Jul 2025 – Present

Las Vegas, NV | Python, SQL, Airflow, PySpark, Snowflake, Trino, Kubernetes

- Independently designed and built a metadata-generator **Airflow** DAG that turns structured inputs, column definitions, unique identifiers, and file-pattern regex, into the **SQL** needed to populate the pipeline's core metadata tables, cutting new-feed setup from about an hour of manual copy-paste-and-edit to a few minutes.
- Deployed 20+ **Snowflake** ingestion pipelines built on this framework, automating data movement from Oracle and SSIS-based legacy systems and cutting manual validation tickets by ~40%.
- Migrated 20+ existing feeds and tables from core- to biz-facing **Snowflake** schemas at the request of the Marketing team, rebuilding SQL views, schema definitions, and target mappings to support their reporting needs.
- Root-caused a recurring **date-parsing failure** affecting multiple Snowflake ingestion pipelines and permanently resolved it at the source, eliminating a class of failures that had required manual reruns.
- Independently troubleshoot pipeline issues spanning Airflow DAG logic, **PySpark jobs on Kubernetes**, Trino SQL views, and Snowflake schema drift.
- Build data quality checks with Great Expectations at each pipeline zone, including column and delimiter validation on incoming files and datatype expectations for types beyond Spark's native varchar/timestamp support, catching malformed business-uploaded files and mis-typed values before they reach downstream tables.

Data Analyst — Advertising Checking Bureau

Sep 2023 – May 2025

Tempe, AZ | SQL, Excel

- Reviewed and validated approximately 120 claims daily using **SQL**, identifying discrepancies that affected reporting accuracy across operational datasets.
- Automated recurring Excel-based reporting workflows, cutting report preparation time from about an hour to under 15 minutes.

EDUCATION

University of California, San Diego — M.S., Data Science

Expected Apr 2027

GPA: 4.0 | Halicioğlu Data Science Scholarship

Arizona State University, Polytechnic — B.S., Information Technology

Aug 2025

GPA: 4.0 | New American Scholarship for Academic Excellence | Dean's List (6 semesters)

Certifications: AWS Academy Cloud Foundations · SQL for Data Analysts and Data Scientists · Tableau for Data Science

TECHNICAL SKILLS

- Languages:** Python (Pandas, NumPy), SQL
- Data Engineering:** Apache Airflow, PySpark, Snowflake, Trino, Kubernetes, Great Expectations, SSIS, ETL/ELT, pipeline orchestration, data validation
- Databases:** Oracle, SQL Server, PostgreSQL, MySQL
- Tools:** DBeaver, GitLab, AWS, Excel, Power BI, Tableau

PROJECTS

Weather ETL Pipeline — Python, GitHub Actions, SQLite

[[GitHub](#), 2026]

- Built an automated ETL pipeline that extracts live weather data from a public API, transforms it into typed records, and loads them into a SQLite database on an hourly schedule via GitHub Actions.
- Designed the pipeline to fail loudly rather than silently — unhandled errors exit with a non-zero status, so a broken run shows up as a failed job instead of a silently missed load.